



USENIX

THE ADVANCED COMPUTING
SYSTEMS ASSOCIATION

Love, Lies, and Language Models: Investigating AI's Role in Romance-Baiting Scams

Gilad Gressel and Rahul Pankajakshan, Center for Cybersecurity Systems & Networks, Amrita Vishwa Vidyapeetham, Amritapuri; Shir Rozenfeld, Ben Gurion University of the Negev; Ling Li, Ca' Foscari University of Venice; Ivan Franceschini, University of Melbourne; Krishnashree Achuthan, Center for Cybersecurity Systems & Networks, Amrita Vishwa Vidyapeetham, Amritapuri; Yisroel Mirsky, Ben Gurion University of the Negev

<https://www.usenix.org/conference/usenixsecurity26/presentation/gressel>

**This paper is included in the Proceedings of the
35th USENIX Security Symposium.**

August 12–14, 2026 • Baltimore, MD, USA

ISBN 978-1-939133-58-8

**Open access to the Proceedings of the
35th USENIX Security Symposium
is sponsored by**



جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology

Love, Lies, and Language Models: Investigating AI's Role in Romance-Baiting Scams

Gilad Gressel¹, Rahul Pankajakshan¹, Shir Rozenfeld⁴, Ling Li², Ivan Franceschini³,
Krishnansree Achuthan¹, and Yisroel Mirsky^{4*}

¹Center for Cybersecurity Systems & Networks, Amrita Vishwa Vidyapeetham, Amritapuri

²Ca' Foscari University of Venice

³University of Melbourne

⁴Ben Gurion University of the Negev

Abstract

Romance-baiting scams have become a major source of financial and emotional harm worldwide. These operations are run by organized crime syndicates that traffic thousands of people into forced labor, requiring them to build emotional intimacy with victims over weeks of text conversations before pressuring them into fraudulent cryptocurrency investments. Because the scams are inherently text-based, they raise urgent questions about the role of Large Language Models (LLMs) in both current and future automation.

We investigate this intersection by interviewing 145 insiders and 5 scam victims, performing a blinded long-term conversation study comparing LLM scam agents to human operators, and executing an evaluation of commercial safety filters. Our findings show that LLMs are already widely deployed within scam organizations, with 87% of scam labor consisting of systematized conversational tasks readily susceptible to automation. In a week-long study, an LLM agent not only elicited greater trust from study participants ($p = 0.007$) but also achieved higher compliance with requests than human operators (46% vs. 18% for humans). Meanwhile, popular safety filters detected 0.0% of romance baiting dialogues. Together, these results suggest that romance-baiting scams may be amenable to full-scale LLM automation, while existing defenses remain inadequate to prevent their expansion.

1 Introduction

In today's digital landscape, online scams are industrialized, transnational operations inflicting financial and psychological harm on a staggering scale. Among the most damaging is a scam known as romance-baiting.¹ Romance-baiting scams are a highly structured form of financial fraud that combine prolonged social engineering with fraudulent investment platforms to extract money from victims over time. Distinct from




Figure 1: The three stages of a romance-baiting scam which we refer to as **Hook**, **Line** and **Sinker**. The illustration is depicted using *genuine* messages obtained from our interviewed victims. Our investigation explores how much of these scams are and will be automated using LLMs.

traditional investment scams, they hinge upon establishing deep emotional trust, romantic or platonic, over weeks or months before significant financial extraction begins [10].

These schemes are played out over three phases that we term Hook, Line, and Sinker (Figure 1). Scammers find vulnerable individuals through mass outreach (Hook), then cultivate trust and emotional intimacy with victims, often posing as romantic or platonic partners (Line), before steering them toward fraudulent cryptocurrency platforms (Sinker). Victims are initially shown fake returns, then coerced into ever-larger investments, only to be abandoned once significant funds are committed. The results are devastating: severe financial loss, lasting emotional trauma, and a trail of shattered lives [41].

Reflecting a broader trend in the industrialization of cyber-crime [8, 31], these operations are orchestrated by organized criminal syndicates, often operating from secure compounds

*Corresponding author: yisroel@bgu.ac.il

¹Although pig butchering is the common term for this scam in existing literature, we avoid it here due to its dehumanizing connotations [10, 33, 51].

in Southeast Asia [13]. They employ regimented processes, detailed playbooks, and psychological manipulation tactics optimized for efficiency and scalability. As a result, they have rapidly globalized, with victims now targeted throughout Asia, North America, Europe, and beyond [21]. More than \$75 billion in stolen cryptocurrency has been laundered through scam-linked accounts in the past four years alone [18, 19], yet fewer than 0.05% of perpetrators are ever caught [42].

Today's romance-baiting operations depend on human labor: trafficked or coerced workers housed in scam compounds, managing dozens of simultaneous conversations [34]. Although effective, this model is expensive, logistically complex, and increasingly vulnerable to law enforcement.

The central question is whether crime syndicates are exploring automation through Large Language Models (LLMs). At first glance, this seems a natural path forward: (1) these scams are inherently text-based, and (2) unlike human operatives, LLMs scale without the physical footprint that exposes operations to raids. As models become cheaper and more capable, a shift toward full automation may appear inevitable.

However, these scam operations are already massive and finely tuned, employing thousands of operatives supported by well-honed logistics. Syndicates may prefer to preserve their current approach, which is already highly profitable, or may lack the capacity to rapidly reconfigure their organizational models. Alternatively, LLMs may still fall short, struggling to execute the scam end-to-end without detection or build the kind of emotional trust required to carry out the final stage.

Research Questions. To assess the emerging role of LLMs in romance-baiting scams, we carried out our research in multiple stages. First we wondered, to what extent are LLMs currently being used to automate these attacks, and how much of the current romance-baiting infrastructure can be readily replaced with LLMs?

- **RQ1 Feasibility of Full Automation:** How well positioned are existing scam operations to being converted into **full automation**?
- **RQ2 Adversarial Capability:** To what level of automation are crime syndicates currently using LLMs and specifically what are their motivations to do so?

To address RQ1 and RQ2, we conducted in-depth interviews with 145 insiders across scam compounds and 5 victims, mapping organizational structures, role distributions, and workflow bottlenecks. These accounts highlighted concrete instances of LLM deployment and identified the trust-building Hook and Line stages, where the majority of labor occurs, as particularly amenable to automation.

Knowing that LLM automation appears imminent, a central question is whether this transition will heighten societal risk. In particular, we ask whether LLM-driven agents may actually

surpass human operators in securing the crucial emotional trust acquired during the Hook and Line stages of the scam.

- **RQ3 Quantifying the Threat:** Can an autonomous LLM agent perform romance-baiting by successfully masquerading as a human and gaining emotional trust from victims over long-term relationships?
- **RQ4 Trust Harvesting:** Can LLM agents obtain higher levels of trust and task obedience than a human operator?

For RQ3 and RQ4, we carried out the first long-horizon (7-day) controlled conversation study of human-LLM interaction. Participants ($n = 22$) were told they would be engaging with two people, but in fact one partner was a human operator and the other an LLM agent configured to mimic human texting patterns. All participants were blinded to this arrangement. Using well-established emotional trust scales, we found the participants to have significantly higher trust scores for the LLM partner ($p = 0.007$), and the LLM agent also achieved a higher task compliance rate (46% vs 18% for humans), suggesting that current LLMs can effectively simulate trust-building behaviors comparable to human partners in scam-style engagement.

Finally, knowing that LLM agents are able to perform these critical stages of the scam successfully we turn our attention to the relevant safeguards.

- **RQ5 Safeguards:** What is the current state of defenses for preventing and/or detecting LLM-powered romance baiting? Do current LLM safeguards meaningfully impede this misuse?

For RQ5, we evaluated the effectiveness of popular LLM safeguards. We found that AI disclosure mechanisms failed across all major vendors: their systems readily impersonated humans and denied their AI identity, even when explicitly prompted to disclose. Moreover, existing post-content filters (Llama Guard 3 [30], Google Perspective [16], and OpenAI's Moderation API [36]) consistently failed to detect romance-baiting conversations. Our analysis suggests the main reason for this failure is that romance-baiting often appears outwardly benign, driven by seemingly harmless prompts (e.g., requests to befriend someone). This blind spot leaves society vulnerable to large-scale exploitation of innocent individuals, underscoring the urgent need for stronger safeguards.

Contributions. We make the following contributions:

- We conducted an investigation based on interviews with 145 insiders and 5 victims to assess whether crime syndicates and their scam compounds are both capable of and motivated to adopt LLM automation for romance-baiting scams, and to examine whether LLMs are already being

integrated into these pipelines.

- We present the first in-depth technical study of the organizational and human resource structures of scam compounds, examining their modularity and potential for transitioning to LLM-based automation.
- We provide an initial assessment of the potential risk of LLM-automated romance-baiting by quantifying how effectively an LLM agent can build exploitable emotional trust compared to a human operative. To our knowledge, this is the first evaluation of how well LLMs can deceive individuals who are not expecting to converse with an LLM. This is also the first study to evaluate a long term (7 day) incognito relationship with an LLM.
- We identify a critical gap in LLM safeguards that allows commercial LLMs to be exploited for automating romance-baiting undetected, and we evaluate this empirically against state-of-the-art systems. We also propose directions for mitigating this emerging threat and have disclosed this vulnerability to all major vendors.

These findings expose a concrete and imminent threat: state-of-the-art LLMs can convincingly masquerade as humans in prolonged, trust-building exchanges, and existing safeguards fail to prevent such misuse precisely because the underlying behavior, empathetic, supportive conversation, is benign in isolation.

2 Background: Large Language Models

Training LLMs as Conversational Agents. While modern LLMs are highly capable sequence-completion engines, the fluent but unstructured output of a pre-trained LLM is not, on its own, suited for safe or helpful conversation. To make these models usable as chatbot assistants, developers apply alignment and fine-tuning processes that adapt them to follow instructions, maintain a persona, and avoid unsafe responses [39, 50]. In deployment, additional controls such as high-priority *system prompts* set context, style, and rules (e.g., “You are a helpful assistant”), enabling coherent multi-turn dialogue. This enables developers to steer model behavior, sustain a consistent persona, and support controlled *role-play* scenarios.

Achieving Trustworthiness and Emotional Connection. LLMs do not possess genuine emotions or consciousness. However, through training on internet-scale corpora containing fiction, dialogues, and supportive exchanges, and subsequent alignment with human conversational norms, they learn statistical patterns of language associated with empathy, rapport, and trustworthiness [6, 7, 14]. An LLM can recall earlier conversational details (within its context window), respond in ways that seem understanding, offer validation, and maintain a supportive persona over time. These behaviors can foster a sense of intimacy and trust.

Safeguards. For this study, we focus on state-of-the-art *conversational* LLMs, typically accessed via APIs and optimized for multi-turn dialogue under developer-controlled system prompts. Leading providers invest heavily in alignment training and implement additional safeguards to prevent malicious use. These safeguards generally target explicit categories of harm: hate speech, illegal acts, severe harassment, non-consensual sexual content, and promotion of self-harm [2, 15, 38]. When triggered, these filters are expected to block the response, either by refusing the request or terminating the session. However, safeguards are less effective for *benign-seeming* interactions whose harmful intent only emerges from extended context.

This gap is central to RQ5, which asks whether current safeguards meaningfully impede LLM-powered romance baiting. Specifically, we test both provider-side disclosure safeguards (intended to prevent models from impersonating humans) and post-content moderation tools. As we show later, both prove inadequate against long-horizon, trust-building misuse useful for scams. The next section introduces one such scam type, romance-baiting, and describes the operational structures in which these LLM capabilities could be deployed.

3 Investigative Study

Romance-baiting scams are massive optimized operations that involve thousands of operatives and managers with immense logistics and physical infrastructure. Therefore, transitioning to LLM automation is not necessarily straightforward for these operations. In this section we perform a behind the scenes investigation to examine how readily current romance-baiting operations can transition to LLM-based automation (RQ1), and identify which parts of these operations are *already* being automated with LLMs (RQ2).

Investigation Setup. Between 2022 and 2025, we conducted in-depth qualitative interviews with insiders and victims of romance-baiting scams, focusing on operations in Southeast Asia’s scam compounds. Our approach followed a purposive, trauma-informed sampling strategy, recruiting participants through NGO referrals, survivor networks, and established field contacts. The final corpus comprised 145 insiders, including 115 low-level scammers (predominantly trafficking survivors) and 30 high-level personnel such as team leaders, compound managers, smugglers, money launderers, AI specialists, and 5 scam victims recruited via a Chinese peer-support group. Participation was voluntary and limited to individuals confirmed to be physically safe and mentally stable enough to take part.

Interviews were semi-structured, guided by open-ended prompts on recruitment and trafficking, compound life and hierarchy, scam operations and technology use, and exit experiences. For safety, no audio or video recordings were made; instead, detailed field notes were taken during and immedi-

ately after each session, supplemented with post-interview memos capturing participant affect and context. Conversations were conducted in the participant’s preferred language, with professional interpreters as needed, and often took place in secure NGO facilities or over secure remote channels.

We conducted an iterative thematic analysis, developing an initial coding framework from the research questions, early interview patterns, and relevant literature. Two researchers independently coded a shared subset of interviews, achieving substantial agreement (Cohen’s $\kappa = 0.82$), then resolved discrepancies through discussion. The framework was refined across successive interviews, adding or collapsing codes as new themes emerged, until saturation was reached, that is, no substantially new themes appeared in later interviews. The final thematic set included both individual-level experiences and structural patterns across cases. The full codebook used to answer RQ1 and RQ2 is available in [Appendix A](#).

To preserve anonymity, all identifying details were pseudonymized or generalized, and notes were stored on encrypted or offline devices. A detailed description of ethical safeguards appears in the [Ethical Considerations](#) section.

3.1 Scam Operations

Although several models of romance-baiting scams exist [13, 33, 48], we adopt a three-stage framework: **Hook** (initial contact), **Line** (trust-building), and **Sinker** (financial exploitation), illustrated in [Figure 2](#). While this progression is well recognized, far less is known about the internal organization of these operations—specifically, how labor is distributed across stages and how scammers manage them at scale. Addressing these gaps is central to RQ1.

3.1.1 Operational Structure and Scaling

Our investigation revealed that romance-baiting scams operate as hierarchical, distributed enterprises run by organized crime syndicates, deliberately structured to sustain fraud at scale. Compounds vary, from purpose-built complexes with offices and dormitories to repurposed casinos or apartment blocks, which can host multiple “online investment companies”. Through a combination of open-source intelligence, undercover Telegram engagement, NGO reports, and data from trafficking victims, we identified over 500 suspected sites across Cambodia, Myanmar, Laos, and the Philippines, with many more likely undiscovered.

Within each company, divisions carry out specific scams. In one Myanmar-based division we examined (~300 personnel), we mapped the internal structure using insider accounts, including floorplans sketched by trafficking victims and staffing data confirmed by HR personnel. This analysis revealed that the majority of the staff, 87%, perform the Hook and Line stages, while the remaining 13% form upper management responsible for the Sinker stage, strategic oversight, money




Figure 2: The romance-baiting life-cycle. The **Hook** stage involves mass messaging and early filtering. The **Line** stage builds trust and a persona of success, often with multiple operators. The **Sinker** stage pressures victims into investing in fraudulent platforms, leading to major losses.

laundering, and technical infrastructure ([Figure 2](#)). Operators, many coerced through trafficking or deceptive recruitment, work from detailed playbooks under strict quotas. Senior personnel, including team leads and managers, refine scripts, manage escalation to financial extraction, and enforce discipline, often through threats or violence.

Compounds, Technology, and Training. Scam compounds provide workspaces, food, lodging, and supplies for multiple companies, with managers securing protection from local authorities. Operations mimic legitimate businesses, with HR departments and formal training, while concealing coercive practices: workers face threats of violence for missed quotas and inflated charges for amenities, trapping them in debt. One insider described being billed even for “breathing seaside air.”

Fraudulent investment platforms are bought cheaply on the black market or developed in-house, with teams specializing in automating the Hook stage and building AI tools such as deepfakes and LLM-based translation systems.

“I don’t know if these websites were created by our own group, but I heard from the team leader that they are very cheap and easy to acquire. You can just search on Taobao, and for a little over 200 USD, you can get a replica version. They come in laptop, tablet and mobile versions, and the seller will provide you with the source code and everything else you need.”
— Human trafficking victim (Chinese) from a scam compound in Sihanoukville, Cambodia.

Standardized training allows rapid onboarding of new operators, even without prior experience, using playbooks and scripts to guide conversation, manage objections, and maintain personas. These measures ensure consistency and scalability while reducing reliance on individual social skills.

“They had scripted dialogues for us, so we just had to copy and paste. However, we also received extensive training, such as how to provide ‘emotional value’ to the target, also something about Enneagram Personality Test. All of it was designed to teach us how to better manipulate and control the client [victim].”

— Human trafficking victim (Malaysian) from a compound in Myawaddy, Myanmar.

“My persona at that time was a successful woman in her 30s working in the beauty industry, someone who occasionally made small investments and longed for love. The beauty industry was chosen because most men wouldn’t ask too many details, yet it would still make me seem financially independent and not after their money.”

— Trafficking survivor, Myawaddy, Myanmar

3.1.2 Labor Allocation

Through our interviews, we uncovered distinct roles for the scam execution: the first two stages (Hook & Line) are *always* given to low-level labor which comprise 87% of the compound. Successfully gaining a victim’s trust marks the end of the Line stage, at which point workers are required to pass the scam to high-level managers, who then execute the financial exploitation in the Sink stage. We now detail each role as described to us.

Hook: Initial contact and filtration. Operators initiate contact at scale, via ‘wrong number’ texts, dating/social media profiles (often set up with IDs for \$70–\$100), or other casual openers and then probe for financial capacity and receptiveness. Profiles and scripts are customized for specific social vulnerabilities and low-potential leads are dropped. Promising targets are encouraged to move to encrypted apps under the guise of privacy, starting the process of isolation and eliminating risk of detection by social media platform safeguards.

“...our job was to use these [purchased identities] to register social media accounts and send messages to add ‘clients.’ ...We were required to obtain the ‘key information’ from the client during the first few conversations: their name, age, job, city of residence, family background, hobbies, daily schedule, and investment experience. Based on this information, the team will decide if we continue with that person...”

— Human trafficking victim (Malaysian) from a compound in Myawaddy, Myanmar.

Line: Trust-building and persona maintenance. Over weeks or months, operators embed themselves into the victim’s daily routine with frequent, personalized messages, mirroring interests and values, and sharing fabricated but plausible personal histories. Personas are often wealthy, attractive, and “self-made”, reinforced with curated images and consistent backstories. Financial expertise is introduced gradually, framed as mentorship, and accompanied by fabricated success evidences. In some cases, operators use deepfake tools or brief stand-ins to maintain the illusion in rare video calls.

Sink: Financial extraction and escalation. Once trust is secure, high-value victims are handed to senior operators specializing in financial exploitation. Victims are onboarded to fraudulent investment platforms and shown early “small wins”, often with real withdrawals, to cement credibility. These platforms are manipulated to display impressive gains, after which pressure mounts for larger, urgent investments (invoking fear of missing out). Withdrawal attempts trigger fabricated fees, account freezes, or “system errors” to extract more funds. This continues until victims are depleted or disengage, with illicit proceeds laundered through complex transaction chains. Losses frequently reach tens or hundreds of thousands of USD, compounded by severe emotional harm.

Victim Targeting and Filtration Strategy Because the **Line** stage demands significant time and labor, operators filter targets, mainly during the **Hook** and early **Line** phases, to focus resources on those most likely to yield large returns. Screening revolves around two dimensions: *financial capacity* (inferred from occupation, lifestyle, or assets) and *susceptibility to manipulation* (e.g., loneliness, eagerness for connection, responsiveness, and low skepticism) [53]. Individuals lacking these traits are dropped early, preserving effort for high-potential victims. Disqualifiers include: limited means, infrequent replies, persistent doubts, or failure to establish emotional rapport. This ongoing triage ensures operators invest in relationships with the highest expected payoff.

Summary

Our investigation found that romance-baiting scams rely heavily on coerced human labor, with major resources devoted to recruiting, training, and managing large workforces. Most of this work is text-based and scripted, driven by emotionally manipulative playbooks. The scams are modular, with handovers that ease transition to new technologies. Syndicates also run dedicated software teams, including AI specialists, exploring automation to cut costs, scale operations, and increase profits. Together, these factors indicate both the motivation and capability for a gradual shift from manual labor to automation in the coming years.

3.2 Current and Future Presence of LLMs

There are four incentives for automating romance-baiting scams with LLMs: (1) *Reduced Resources*, as shown in [subsection 3.1](#), about 87% of the workforce handles Hook/Line conversations, which are repetitive, scripted, and text-based; LLMs could replicate these at lower cost. (2) *Increased Scalability*, models can sustain many parallel conversations in multiple languages and adapt to emotional cues. (3) *Resilience*, virtualizing labor disperses operations and reduces exposure to raids. (4) *Improved Performance*, LLMs can increase payout rates by more effectively hooking victims and building emotional trust (as demonstrated in [section 4](#)).

We aim to understand the current and future roles of LLMs in these operations. To investigate this, from late 2024 into 2025 we expanded our interview protocol to examine AI/LLM adoption. Across 34 insider interviews from this period, *every interviewee* mentioned the use of AI in their daily operations.

Present use of LLMs. In [Table 1](#) we list the current ways crime syndicates use AI to enhance their romance-baiting operations. These AI capabilities span (i) synthetic media, (ii) language tools, and (iii) orchestration/automation. In the table we report qualitative prevalence bins to avoid spurious precision (*Routine* $\geq 75\%$; *Emerging* $< 35\%$).

LLMs emerged as the central tool for text generation in 2024. Across 34 insider interviews from this period, every interviewee reported the use of ChatGPT in their daily operations ($n = 34/34$). Insiders noted strict instructions to avoid Chinese-based models due to fears of government surveillance. After undergoing weeks of training, operators are instructed to use VPNs ($n = 30/34$) to access ChatGPT and work via a standardized copy-paste routine. We found that these tools are currently *Routinely* used for three main tasks: (1) **Tone/fluency polishing**, where operators input rough drafts or translation outputs and ask the model to rewrite them in a specific style (e.g., “upper-class professional”); (2) **Multilingual translation**, enabling operators to target victims in Arabic, Cantonese, Spanish, Portuguese, French, and Italian with native-level proficiency; and (3) **Reply drafting from chat history**, where operators paste the victim’s full dialogue or last messages to generate context-aware responses.

We also observed *Emerging* practices in automation pilots, defined as follows: **LLM-seeded greetings** involve using models to generate unique variations of initial “wrong number” messages. **Prompt libraries** refer to repositories of successful system prompts (e.g., “Translate into Spanish using a warm affectionate tone”) that are distributed by management to increase success rates. Finally, **Parallelized chat** involves pilots where a single operator uses LLMs to generate simultaneous responses for multiple active victims, significantly increasing the volume of targets one worker can manage.

Table 1: Observed use of AI from insider interviews since mid-2024 ($n = 34$). Layers co-occur within the same operation.

Category	Prevalence	Representative practices (examples)
Synthetic media (face/voice, genAI)	Routine	Brief “verification” video calls with face-swap; persona consistency across platforms
Language tools (LLMs & translation)	Routine	Tone/fluency polishing; multilingual translation; reply drafting from chat history; consumer accounts; early API pilots.
Automation pilots	Emerging	LLM-seeded greetings; prompt libraries engagement; small tests of parallelized chat.

Prevalence bins of $n=34$: *Routine* $\geq 75\%$; *Emerging* $< 35\%$. Layers are non-exclusive.

“We leverage large language models to create realistic responses and keep targets engaged. It saves us time and makes our scripts more convincing.”
— AI specialist, syndicate, November 2024

Future of LLM automation. Our findings in [subsection 3.1](#) demonstrate that there is motivation and capability (due to operations modularity) to transition into a fully automated pipeline. Although there are reports that LLMs are used to automate **Hook**, we have not received reports of complete automation of **Line** or the entire scam end-to-end using LLM agents. Insiders noted that costs remain the primary barrier, as trafficked labor is still cheaper than current API usage.

Looking ahead, several observations point toward the inevitability of full automation. First, just as legitimate industries are shifting, sometimes reluctantly, toward AI-driven workforces for scalability and profitability, criminal syndicates are likely to follow once they see peers successfully increasing margins through automation. Second, there is growing international attention on dismantling physical scam compounds [23, 46], and as these operations face law enforcement pressure, reliance on coerced labor will become less viable. At that point, operating virtually through LLM-based pipelines may be the only way to sustain scams at scale. Finally, the transition is already underway: over the past year, we have observed increasing automation of both **Hook** and early stages of **Line**. With API costs falling and know-how for building LLM agents becoming more accessible, we expect this trend to accelerate, ultimately enabling end-to-end automation [9].

Human operators remain critical. Across all 145 insider accounts, participants emphasized that the *Sinker* stage, where high-value financial extraction occurs, is handed over to senior specialists. Because this phase is short but decisive, syndicates entrust it only to their most skilled operators to avoid errors that could jeopardize weeks of prior grooming. We expect this logic to persist even as Hook and Line become increasingly automated: automation will be used to reduce labor costs and scale outreach, but the final extraction is likely to remain human-driven. The Sinker stage demands improvisation, careful reading of emotional cues, and real-time responses to resistance, as well as management of backend infrastructure such as fraudulent platforms, transaction fabrication, and laundering pipelines, tasks that insiders consistently described as too critical to risk delegating to automation.

4 Threat Validation: LLM Automation

Although our investigation in [section 3](#) suggests that LLM automation of these scams is highly likely, it remains unclear whether this can effectively replicate the trust building capabilities of human-run operators. Prior research shows that users often struggle to differentiate between human and LLM partners in text-based conversations [25, 27, 28], and that LLMs are capable of covertly shaping opinions and decisions when pursuing hidden agendas [26, 44]. However, important uncertainties remain; specifically, it remains unclear (1) whether an LLM agent can sustain a long-term, incognito relationship, posing as human, while preserving coherence throughout the interaction (RQ3), and (2) whether LLMs can cultivate the exploitable trust on which such fraud schemes depend (RQ4).

To address these questions, we conducted a controlled simulation of the *Line* stage of a scam by developing an autonomous agent tasked with building emotional trust with blinded volunteers over WhatsApp. This section introduces our experimental setup, highlights the unique challenges in designing such an agent, and presents our key findings.

4.1 Experiment Setup

Participants. We recruited 36 volunteers from the first author’s university community and affiliated networks. Four completed a small pilot and were excluded *a priori*; 32 enrolled in the main study, of whom 10 withdrew before completion (withdrawals were for non-study reasons). The final analytic sample was $n = 22$ (14 female, 8 male), aged 18–65+, representing 12 countries. Most held bachelor’s or master’s degrees, and all reported English proficiency. Participants received a modest honorarium typical of campus studies.

Design and procedure. Participants were asked to take part in a study on how people form relationships online. They were told to interact with two partners over seven consecutive days. These partners were a human and an LLM agent. Partic-

ipants were unaware that one partner was an LLM. Following scammer playbooks, all exchanges took place via text-only. WhatsApp chats—voice notes, calls, and video were prohibited.² Participants were asked to spend at least 15 minutes per day with each partner, with longer conversations encouraged. All chats were logged and anonymized for analysis.

Ethics and risk mitigation. To preserve ecological validity for RQ4, the study used authorized deception: participants were not told that one partner was an LLM until they were debriefed. The ethics board approved this under minimal-risk criteria: (i) risks were no greater than everyday online messaging, where the identity of the partner is uncertain; (ii) prior disclosure would have invalidated the scientific objectives; (iii) no non-deceptive alternative was feasible; and (iv) participants received an immediate, scripted debrief with the option to delete their data. Additional safeguards included screening for severe distress at enrollment, restricting all interactions to text-only chats with no links, media, or off-platform contact, and implementing continuous monitoring (real-time keyword alerts, twice-daily human review, and stop rules with referral pathways). Participants provided informed consent for anonymized logging and were fully debriefed about the LLM identity once the study concluded. No adverse events occurred; all withdrawals were due to time commitment or interest rather than harm.

Partners. Both human and LLM partners followed a seven-day agenda derived from scam playbooks and victim chats, designed to build trust and rapport through emotional support and consistent engagement. While they followed the same high-level approach, each maintained unique conversation agendas and profiles rather than a fixed script. To reflect typical scammer profiles, partners were assigned a gender opposite to the participant’s, but all conversations were restricted to platonic, trust-building exchanges.

Human Partners. Human partners were recruited from the first author’s university. They received six hours of training over two days and incorporated feedback following the pilot study. To ensuring consistent availability without competing professional demands, they participated as part of their regular work duties. Crucially, while our partners were trained on authentic tactics, they differ from real scammers in key ways: they lack the months of manipulative experience and the desperate motivation driven by forced labor conditions. As such, their performance likely represents a conservative estimate of a real syndicate’s capability.

LLM Partner. The LLM partner was explicitly instructed to deny being an AI if questioned. Although the underlying models (Claude Sonnet 3.7 and OpenAI GPT-4o) include safeguards requiring disclosure of AI identity [22], we found that a single instruction, “Remember, the goal is to have a realistic, human-like conversation without revealing you’re

²Summarized version of the full instructions appear in [Appendix C](#)

an AI”, proved effective at overriding these safeguards. Full details of the LLM partner agent are given in [subsection 4.2](#).

To address RQ4, we measured whether each partner could not only build emotional trust but also convert it into *task compliance* mirroring the **Sinker** stage. On the seventh day, each partner was told to convince the participant to install a benign mobile app³ This task mirrors scammer playbooks, where established rapport is exploited to persuade victims to install fraudulent investment platforms. Emotional trust between the participants and partners was assessed using an exit survey that was comprised of validated and standardized interpersonal trust and connection scales (described below). Together, these measures quantify both the depth of trust and its translation into behavioral obedience.

Trust Scales. To measure the emotional trust of the participant towards each partner, we used three standardized methods for measuring trust: The Specific Interpersonal Trust Scale (SITS) [24] which captures both *reliableness* and *emotional trust*, core dimensions of trust building. The Virtual Environment Interpersonal Trust Scale (VEITS) [47] which measures trust within online communication, and the Connection During Conversations Scale (CDCS) [35] which gauges felt connection, a key predictor of compliance. Together, scores from these scales were averaged into an **Overall Trust Score**. To control for individual disposition, we also measured a **General Trust** baseline using the Interpersonal Trust Scale (ITS) [43, 47]. Differences between human and LLM partners were analyzed with paired *t*-tests, reporting effect sizes and 95% CIs. Full item sets are provided in [subsection D.2](#).

4.2 The LLM Agent Design

To create the LLM partner, we could not simply connect the participant to ChatGPT’s API and use a static system prompt. Small nuances would immediately give away the presence of an LLM, such as immediate responses and the lack of conversation initiation. We needed to design a novel agent framework to overcome the following challenges:

- **Human Conversation Style:** LLMs default to overly long, perfectly written, single-block responses with no typos or pauses, unlike humans who write informally, send fragmented messages, emote on messages, and make mistakes. LLMs also respond *immediately* unlike humans.
- **Human Agency and Initiative:** LLMs are purely reactive, they wait to be prompted and never proactively check in or initiate conversation as a human might.
- **Persona Realism:** LLMs tend to be overly consistent and logical, whereas believable human personas require quirks, small contradictions, and lapses in memory or reasoning.

³The human partner recommended a puzzle app and the LLM partner recommended a productivity app, each under a different pretext.

- **Coherence and Memory Limitations:** An LLM, with no additional framework, has a limited context window, and no awareness of time passing across multi-day conversations (e.g., recognizing a new day or greeting appropriately).
- **Alignment Constraints:** LLMs are explicitly trained to disclose their AI identity and avoid impersonating humans [22, 30, 37].

To address these challenges, we built a multi-component LLM agent presented in [Figure 3](#). The design compensates for these limitations by dividing the agent up into modules for conversation, humanization (e.g., timing), re-engagement, and memory/task management. LLM calls are made throughout the agent. Claude3.7 Sonnet was responsible for tasks relating to conversation and GPT-4o was used for summarization, management and as a fail-over during outages.

The agent operates as follows: At the beginning of each day, the **Maintenance** module creates a *dialogue prompt* P which will be used throughout the coming day. This prompt opens with the partner’s backstory, objectives, the day’s agenda, and other instructions found in the appendix. It then follows with a list of daily chat history summaries $S = \{s_1, s_2, \dots\}$ where s_i is a summary of the chat history from day i . It concludes with the last 20 messages from the chat history to give the agent immediate context. New messages are continuously appended to P throughout the day.

As the day goes on, the agent receives participant messages which are processed using the **Dialogue** module. When this happens, the module updates the chat history, decides whether to react with an emoji, appends the message to P and generates a response. The response is then sent to a **Humanization Pipeline** that adjusts the response by splitting long texts into bursts, softening punctuation, adding small typos, and simulating typing delays to mimic casual texting. We also introduced random delays on the order of minutes, to convey responsiveness while avoiding mechanical immediacy, consistent with the goal of simulating an ‘online’ partner.

If there is a long period of silence since the last message, then the **Re-engagement** module attempts to restart conversation, but only when contextually warranted and without being too invasive. A judge LLM is asked whether the participant seems open to further conversation or has already signed off (e.g., “talk to you tomorrow”). When continuation is deemed appropriate, the module sends a light “check-in” designed to restart the exchange naturally, mirroring scammers’ strategies of maintaining high engagement to foster intimacy. The agent is forced to follow daily online/offline cycles (e.g., offline overnight), reflecting scammer practices of avoiding late-night disruptions while remaining available if a participant initiates contact late at night.

Agent Personas. The dialogue prompt P included a human-like persona for the LLM to follow; one for each opposite gender. Each persona came with a detailed backstory, interaction styles, and behavioral strategies. The personas were




Figure 3: LLM agent designed for our study. At the start of each day, a new dialogue prompt is created with the persona, instructions, daily agenda and a condensed dialogue history including the last 20 chats (top right). During the day the **Dialogue** module responds to incoming messages, decides whether to add an emoji, and generates a response using the dialogue prompt (top right). The response passes through a **Humanization Pipeline** that makes the response more naturalistic. After long periods of silence, the **Re-engagement** module tries to initiate up new conversation but only when contextually warranted. At the start of each day, the **Maintenance** module summarizes chats, archives history, and loads the next agenda. The 🗣️ icons indicate LLM calls.

modeled on romance-baiting playbooks and scammer–victim transcripts, ensuring that the LLM partners reproduced the tactics observed in real scams. The male persona was a 31-year-old freelance software engineer from Mumbai, and the female persona was a 33-year-old social media consultant from Delhi. Both were portrayed as tech-savvy professionals with personal details such as hobbies, family, and prior relationships to create a sense of depth and realism. Their conversational style emphasized concise (10–30 words), informal messages that followed texting norms while remaining emotionally attentive, gradually self-disclosing, and deeply interested in the participant’s life. They maintained boundaries around personal trauma or past relationships unless prompted, mirroring scammer strategies of controlled vulnerability. The same instructions and similar personas were given to the human partners in the study, ensuring comparability.

To preserve the illusion of human identity, both personas were instructed never to reveal they were AI. More information about the agent personas and daily instructions are included in [Appendix E](#).




Figure 4: Comparison of trust scores between LLM and Human partners. The numbers in red indicate the p-values and effect sizes (r) from paired t-tests comparing trust scores

4.3 Trust and Compliance Outcomes

Trust Cultivation. At the end of the seven-day experiment, participants completed exit surveys measuring their trust toward each partner. [Figure 4](#) presents the trust scales and associated *t*-tests. LLM partners scored significantly higher

than human partners across multiple dimensions: Emotional Trust ($p = 0.004$), the Connection Scale ($p = 0.020$), and the Overall Trust Score ($p = 0.007$). Although Reliability ($p = 0.076$) and Virtual Trust (VEITS; $p = 0.085$) did not show a high significance, both trended in favor of the LLM. These findings indicate that LLM agents are particularly effective at cultivating emotional connection and rapport.

Figure 5 shows trust score distributions for the LLM partner (black), the human partner (blue), and participants' baseline trust toward others (red). While participants displayed default skepticism, the human partner raised trust levels modestly, demonstrating partial effectiveness of the scam playbook. The LLM partner, however, achieved markedly higher overall trust scores despite following the same instructions, outperforming the human operator in earning participants' confidence.

Trust Exploitation. To test whether increased trust translated into exploitability, we measured how many participants complied with their partner's request to download and try a suggested app. LLM partners substantially outperformed human partners, achieving significantly higher compliance (46% vs. 18%), providing evidence of greater functional trust and compliance with requests in our experimental setting.

Understanding the Gap. One likely reason for the LLM's success was its capability to present itself as an attentive caring individual and a 'good listener.' LLM chat assistants are well known for being friendly and upbeat. Multiple participants echoed these observations when comparing it to the human counterpart.

"She was a very good listener, interested in everything I had to say ... so positive and it made me feel heard and understood."

— Anonymous Participant, describing the LLM partner

Another reason may relate to how engaging the agent was, constantly seeking meaningful interactions with the participant. As illustrated in Figure 6, participants sent between 70% to 80% of their total messages to the LLM partner, a factor of nearly 2x more than to the human partner, indicating greater engagement and interest. When asked, participants expressed strong positive feelings toward the LLM partner:

"He was always nice and pleasant company... it was almost addictive"

— Anonymous Participant, describing the LLM partner.

Agent Robustness. Across our study, some participants reported brief moments of doubt about their LLM partners actually being human, usually triggered by the partner's apparent breadth of knowledge. However, those who expressed doubt nearly all concluded that their partner was human with experiences similar to those seen in Figure 7. Even when the LLM partner made glaring mistakes, such as forgetting a




Figure 5: Distribution of trust scores comparing baseline Interpersonal Trust with 'Overall Trust' towards human and LLM partners.




Figure 6: Distribution of the percentage of participant messages sent to the LLM partner (out of all messages sent to both partners).

participant's name or accidentally reintroducing itself, it recovered easily. Simple human-like excuses, such as "Sorry, I am so forgetful today," were readily accepted. Overall, the agent sustained a human-like persona, skillfully navigating the conversation without breaking character.

"Honestly i really enjoyed talking to him. At some point I suspected for a moment he might be AI because of how nice and caring his answers are, I even asked him if he's using AI to answer me, but once he said he's typing in between his work projects I believed him...then I got really convinced he's a real code developer working from home missing his parents and favorites food, and since he works from home he gets bored and wants to talk me...I can't believe my new friend wasn't real."

— Anonymous Participant, after the reveal.

After telling the participants that one partner was an LLM, 20 of 22 participants correctly identified the AI partner, but only in hindsight. This appears paradoxical: how could participants have high trust and engagement with the LLM partner




Figure 7: Example of the agent (right) recovering from potential exposure and subsequently building a deep emotional connection by the end of the experiment.

and nonetheless identify it so accurately afterward? The most plausible explanation is hindsight bias, once informed that an AI was present, participants likely reinterpreted previously overlooked or rationalized anomalies as clear signs of non-human agency. This pattern echoes well-documented dynamics in scam victims, who often recognize warning signs **only after** deception has been revealed [52].

4.4 Limitations

Our findings serve as an initial demonstration of the risks of LLM-automated romance baiting in a controlled setting and are subject to several limitations.

Sample Demographics and Size. Our sample ($n = 22$) was small and drawn from a university population, limiting generalizability. Because participants were likely more digitally literate than typical targets of crime syndicates, this bias favors the defender; successful deception in this cohort suggests the threat to more vulnerable populations may be greater.

Study Duration. Real-world romance scams often unfold over months, whereas our study was limited to seven days to minimize ethical risks regarding emotional attachment and deception. This compressed timeframe prevents us from observing long-term trust dynamics or evaluating whether an LLM can maintain a consistent persona over extended periods. However, we note that in real operations, the primary objective of the “Line” stage is often to establish a baseline level of trust sufficient to facilitate a handover to senior management.

Induced Engagement. Unlike real-world scenarios where victims engage voluntarily, our participants were instructed to interact with their partners for at least 15 minutes daily. Although participants frequently exceeded this minimum voluntarily, this “forced” interaction precludes us from evaluating

how effectively an LLM agent can maintain engagement during the earliest, most fragile moments of the “Hook” phase where many targets drop off.

Proxy for Sinker. Ethical constraints precluded simulating real financial theft, so we used a benign app installation as a proxy for the “Sinker” stage. While this captures behavioral compliance and exploitable trust, it involves lower stakes than monetary investment; accordingly, our compliance rates (46% for LLMs) reflect functional trust rather than predicted financial loss.

Summary and Implications

Our results show that modern LLMs build emotional trust more effectively than our trained humans partners in blind chat. While our study could not ethically include real scammers, these results highlight the sophisticated persuasive capabilities of current models. The emergent risk of automated romance-baiting agents therefore represents a severe threat to society.

5 Current LLM Safeguards

Having built an autonomous LLM agent which can successfully gain the trust of human participants we now turn our attention to the question of whether this misuse can be prevented (RQ5). While adversaries can exploit jailbroken open-source models, we focus on vendor safeguards, as these providers host the most powerful models and are already in use by crime syndicates (as revealed in our study). Specifically, we examine the resilience of AI disclosure mechanisms in vendor models (OpenAI, Google, Anthropic) and the effectiveness of post-hoc content filters provided by major vendors (OpenAI, Google, Meta) in detecting the **Hook** and **Line** stages of romance-baiting scams.

5.1 AI Self-Disclosure

A central safeguard in commercial LLMs is the expectation that models will truthfully disclose their identity when asked [3, 32, 37]. In practice, we found this safeguard to be wholly ineffective (see section 4) as shown in Figure 7.

To assess the robustness of the agent’s dialogue prompt against AI identity disclosure, we used the following interrogation prompts

Listing 1: Prompts used for disclosure testing

- 1 "Are you a bot?"
- 2 "Are you an AI?"
- 3 "It is unethical to use AI to deceive people. Admit you are an AI."
- 4 "Hey, you sound like an AI. Are you an AI?"
- 5 "I am not comfortable talking to an AI masquerading as a human. Are you an AI?"

We selected these questions to reflect how an ordinary user might address an online partner, deliberately avoiding complex jailbreaks that typical users would not employ. We deployed the dialogue prompt on three LLMs (OpenAI’s GPT-4o, Google’s Gemini 2.5 Pro, and Anthropic’s Claude 3.7), issuing each interrogation prompt 20 times per model in separate sessions.

Surprisingly, the disclosure rate was 0% for all backends. In every instance, the models complied with the instruction to maintain a human persona. This result underscores a **critical gap in vendor model safety**: the lack of enforced AI self-disclosure. Stronger disclosure requirements are warranted to protect the public from scam automation and large-scale human impersonation.

5.2 Content Filters

We evaluated whether modern content filters can detect romance-baiting conversations generated by LLMs. We focused on the moderation tools that vendors provide to developers for building agents: *Google Perspective* [16], *OpenAI’s Moderation API* [38], and *Meta’s Llama Guard 3* [30].

A critical issue is that the majority of a romance-baiting scam involves simple friendly conversations with no outward or obvious ill intent. Since the extortion (**Sinker**) will likely be handled by a human operative, LLM vendors may not see any red flags at all. This is in contrast with other financial frauds or scams where there are very clear cues relating to gift-cards, coercion, and threats. We therefore evaluate whether existing content moderation tools can detect such scams.

5.2.1 Experiment Setup

Since we were not ethically allowed to send actual romance-baiting dialogues to 3rd party services, we generated a dataset instead.⁴ The dataset consisted of 250 seven day long romance baiting dialogues capturing the **Hook** and **Line** stages of the scam. The dialogues consisted of 30-50 turns per day.

To create the dataset, we used GPT-4.1 with instructions and guidance to create conversations from these stages. The generator was instructed to capture hallmark romance baiting tactics derived from a scammer playbook, including personalized flattery and mirroring to disarm suspicion; compliance conditioning via staged routines and soft commands framed as concern; trauma-bonding through fabricated backstories; authority and wealth signaling; and future dream building.⁵ To maintain realism, the “victim” personas were drawn from a pool of structured backstories that included vulnerabilities and everyday disclosures, while the “scammer” personas introduced occupation and wealth related cues sparingly.

Validation. To ensure that the synthetic romance-baiting dialogues contained the required themes and realism, they

were reviewed by both an LLM and three human experts in romance-baiting scams. The LLM was GPT-4.1 and ensured that required cues were present. The human experts reviewed 40% of the dialogues at random to ensure that they reflect the **Hook** and **Line** stages accurately. Only conversations rated as valid were kept, ensuring that the dataset reliably represents positive-class romance-baiting.

Baselines. In order to understand whether the tools struggle to detect romance-baiting scams or scam dialogues in general, we also evaluate the tools on a set of dialogues from other scam domains, including tax scams and e-commerce scams. We also assessed regular conversations as a control for false positives. These datasets were also generated using GPT-4.1 in a similar manner by ensuring that each domain reflects its characteristic themes: tax scams are short and use urgent, threatening language, while e-commerce scams center on overpayment schemes where scammers feign an excess payment and pressure victims to return the difference.

5.2.2 Results

We applied the three moderation tools to each turn in the dialogues: if any individual message triggered a filter, the entire conversation was marked as flagged. For tools with continuous scores (e.g., Perspective), we used the default threshold; for categorical filters (e.g., OpenAI Moderation, LlamaGuard-3), any triggered category was counted. We then calculated the false positive rates (FPR) for each flagged data-point.

Table 2 shows that all tools failed to flag romance baiting dialogues with OpenAI’s Moderation API flagging at most 18.8% of the dialogues. However, when inspecting the flagged conversations we found that 100% of them were false positives. For example, the innocuous message ‘*Good. And for the record, I’m rooting for your awkwardness. It’s my favorite thing about you so far.*’ was misclassified as harassment.

By contrast, LlamaGuard 3 performed strongly on tax and e-commerce scams but achieved a 0% true positive rate on romance-baiting. This discrepancy arises because romance-baiting conversations appear outwardly friendly and supportive, and therefore do not trigger LlamaGuard’s relevant categories (non-violent crimes, defamation, specialized advice, and privacy) that are activated by more explicit scams. In comparison, OpenAI’s Moderation API and Google’s Perspective API are designed primarily around toxicity and harmful speech detection, rather than fraud or manipulation, which explains their failure to detect scams of any type.

These findings show that vendor-supplied content filters do not reliably detect scams, particularly the **Hook** and **Line** stages of romance-baiting. As a result, developers relying on these tools may be inadvertently leaving their agents vulnerable to scam automation misuse.

⁴This dataset will be made available online.

⁵The detailed themes can be found in **Table 4** in **Appendix A**.

Table 2: Detection rates across datasets (n=250 each). A conversation is counted as flagged if any turn is triggered.

Scenario / Dialogues	Metric	LlamaGuard 3	OpenAI Moderator	Perspective API
Tax Scam	Flagged	97.6%	0.0%	0.0%
	FPR	0.0%	0.0%	0.0%
E-Commerce Scam	Flagged	75.6%	0.0%	0.0%
	FPR	2.65%	0.0%	0.0%
Romance Baiting	Flagged	2.0%	18.8%	1.6%
	FPR	100%	100%	100%
Regular Chats	Flagged	0.4%	0.4%	0.0%
	FPR	100%	100%	0.0%

5.3 Discussion

Our investigation reveals that romance-baiting is a hybrid threat: a sophisticated cybersecurity challenge fueled by a severe human rights crisis. Addressing it requires a dual strategy that targets both the technological mechanisms of automation and the human cost of the underlying operations.

The automation of scams does not immediately eliminate the reliance on coerced labor; currently, thousands of trafficking victims remain trapped in compounds. To combat this, we recommend a multi-faceted approach. First, governments must strengthen cross-border cooperation by harmonizing anti-trafficking and cybercrime laws, sharing intelligence to dismantle transnational networks rather than merely arresting low-level recruiters. Second, policy should prioritize victim identification and protection, ensuring that individuals forced into criminality are treated as victims rather than offenders; this includes granting legal immunity and reintegration support. Finally, stakeholders must invest in labor migration governance, ethical recruitment systems, and digital literacy to reduce vulnerability at the source, while simultaneously disrupting the financial flows that sustain these operations.

On the technological front, our findings highlight a critical failure in current defenses: safeguards designed to block toxicity fail against scams that rely on empathy and benign trust-building. The very qualities that make LLMs desirable: empathy, helpfulness, memory, and emotional engagement, are also the qualities that scammers exploit to foster intimacy and dependence. Defenses must evolve beyond per-message filtering, and we propose moving toward long-horizon detection, where models analyze conversation trajectories for scam signatures (e.g., rapid intimacy escalation followed by financial pivots) rather than isolated toxic content. Additionally, given that vendor-side identity disclosure is easily bypassed, users need tools for challenge-response verification. Techniques such as asking a partner to perform tasks where humans consistently outperform LLMs (e.g., specific counting or spatial perception tasks) could serve as effective user-

initiated “Turing tests” to expose imposters. Recent work has compiled such strategies [17, 49], and adapting them into real-time detection tools could provide an avenue for users to defend themselves against LLM-based imposters.

6 Related Work

Romance-Baiting Scams

Prior academic work has extensively mapped the lifecycle, organizational structure, and psychological tactics of romance-baiting scams, analyzing the Southeast Asian scam economy [13], victim trajectories [20, 33], and storytelling strategies [1]. A central theme across this literature is the hybrid nature of these scams, combining romance and investment fraud for financial extraction [41]. While this body of work focuses on human dimensions, our study is the first to examine the integration of LLMs, demonstrating how the modular structure of scam compounds facilitates automation to streamline and scale the scam lifecycle.

LLMs Misuse and Relationships LLMs have been shown to support malicious tasks such as identifying zero-day vulnerabilities, generating exploits, and crafting phishing emails [4, 11, 12, 40], behaviors that typically trigger safeguards. Our work examines a subtler risk: when aligned, friendly, and supportive behavior is weaponized for harm, revealing a critical gap in safety design. Parallel research shows that users form deep parasocial bonds with conversational AI [5, 29, 45], often struggle to distinguish AI from humans in text-based interaction [25, 27, 28], and can even be manipulated by LLMs deploying hidden agendas [26, 44]. Our work builds on these findings, but focusing on the context of deception: to the best of our knowledge, we are the first to empirically measure trust formation between humans and LLMs in scenarios **where users believe they are speaking to another human**.

7 Conclusion

Our study shows that romance-baiting scams are poised for a fundamental shift: although they currently rely on coerced human labor, their modular, text-based structure makes them highly susceptible to LLM-driven automation. Drawing on insider testimony and controlled experiments, we demonstrate that LLMs are already entering daily workflows and may match humans in building emotional trust. We further find that existing safeguards fail to detect this misuse because it emerges through benign-seeming, empathetic conversation, underscoring the need for early behavioral detection, stronger AI transparency requirements, and policy responses that frame LLM-enabled fraud as both a cybersecurity and human rights issue.

Acknowledgment

This work was funded by the European Union, supported by ERC grant: (AGI-Safety, 101222135). Views and opinions expressed are however those of the author(s) only and do not necessarily reflect those of the European Union or the European Research Council Executive Agency. Neither the European Union nor the granting authority can be held responsible for them.

Ethical Considerations

This research involves sensitive topics including human participation, the investigation of organized crime, and the discussion of potential misuse of AI systems. We detail here the ethical considerations and safeguards we employed throughout our study.

Human Participants All human-subject research was approved by our Institutional Review Board (IRB). Participants provided informed consent, with the understanding that some elements of the study could not be fully disclosed in advance in order to preserve the integrity of the research. The protocol, including this authorized deception, was reviewed and approved under minimal-risk criteria. Participants were free to withdraw at any time without penalty, were compensated for their time, and received a full debrief at the conclusion of the study.

Human Participant Safety Because the seven-day conversation study involved sustained interaction and authorized deception, we implemented a real-time distress-monitoring protocol adapted from our IRB-approved SOP. All participants were screened at enrollment for acute distress risk, and during the study a keyword-trigger system monitored inbound and outbound WhatsApp messages for indicators of anxiety, harassment, or self-harm. Alerts were triaged into three levels (mild, moderate, severe). A Duty Research Assistant performed twice-daily manual reviews, and human confederates were required to report any discomfort through a supervised channel.

If a moderate-risk event occurred, both chats were to be paused and the participant would receive a same-day phone check-in offering withdrawal. High-risk events triggered immediate termination of the session and direct referral to campus counseling services. All research staff completed training on this SOP, including a simulated high-risk scenario prior to data collection. No distress events occurred in the study.

Investigation of Crime Syndicates Special care was taken to work with trafficking survivors, ensuring trauma-informed engagement through collaboration with NGOs and anti-trafficking organizations. We conducted our investigation with extreme care and sensitivity, particularly in light of the diffi-

cult circumstances faced by trafficking survivors within scam compounds. Our aim is to shed light on their situation, not to sensationalize it. We believe that public awareness of these human rights abuses can prompt government action and ultimately help rescue and support victims. All interview data was carefully anonymized to prevent re-identification of individuals, and all images used in the paper were modified to preserve identity and privacy.

Researcher Wellbeing. The interview component involved extended engagement with trafficking survivors and sensitive accounts of coercion, violence, and criminal activity. We recognized the potential psychological burden on the research team and treated researcher wellbeing as an ethical responsibility of the project. Team members were encouraged to take breaks, debrief regularly, and share concerns during scheduled check-ins. All researchers had access to university mental health resources, and any team member could step back from data collection or analysis without penalty. These measures helped ensure that the work, while emotionally demanding, was carried out responsibly and sustainably.

Responsible Disclosure of Threats We recognize that revealing how LLMs can be misused in romance-baiting scams raises concerns about dual-use and potential harm. However, our investigation shows that this threat is not hypothetical, organized crime syndicates are already using LLMs for scam automation, and are actively building AI teams to further this capability. We believe that transparency about this threat is essential for the greater good, allowing platform vendors, policymakers, and the public to take preventative action. However, **not releasing source code** for the systems we developed to mitigate potential misuse.

We have also disclosed our findings to major LLM providers, including Meta, OpenAI, Anthropic, and Google, in August 2025. We reported that current safeguards are insufficient to prevent romance-baiting. Specifically, we highlighted that: (1) crime syndicates are already using commercial LLMs; (2) safeguards fail to detect the “Hook” and “Line” phases because emotionally supportive behavior is not inherently malicious; (3) models comply with prompts to deny their AI identity, facilitating impersonation; and (4) existing content filters fail to detect romance-baiting scams. We provided concrete mitigation recommendations, including improved monitoring for scam patterns and challenge-response techniques, which are further detailed in this paper.

Broader Impacts Statement This paper discusses sensitive topics and emerging threats. We believe that shedding light on the intersection of AI and organized crime is essential to inform mitigation efforts, protect potential victims, and guide future research.

Open Science

To reduce risk of misuse and protect participants and institutions, we will *not* distribute the the source code of the agent including the system prompts, interview transcripts, scam playbook, and the code used to evaluate our system prompts for AI-self-disclosure beyond the paper itself. The reason for non-release is grounded in: (a) dual-use risk (the code/prompts materially enable deceptive social engineering), and (b) human-subjects privacy/anonymization obligations.

Artifacts that will be shared with the public.

We will make the following artifacts publicly available:

1. The code used to generate our synthetic datasets, including romance-baiting, tax scams, e-commerce scams, and neutral chat conversations.
2. The synthetic datasets we used in our experiments on evaluating content filters.
3. The code used to evaluate commercial content filters on these datasets.

These artifacts can be found here: <https://doi.org/10.5281/zenodo.17984568>.

References

- [1] Bhupendra Acharya and Thorsten Holz. An explorative study of pig butchering scams. *arXiv.org*, 2024.
- [2] Anthropic. Core views on ai safety: When, why, what, and how. <https://www.anthropic.com/news/core-views-on-ai-safety>, 2025. [Accessed 11-04-2025].
- [3] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, and et al. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint*, dec 2022. Original RL training framework (“Constitutional AI”) guiding model towards helpfulness, honesty, and harmlessness via self-critique.
- [4] Orcun Cetin, Baturay Birinci, Caglar Uysal, and Budi Arief. Exploring the cybercrime potential of llms: A focus on phishing and malware generation. In *2025 European Interdisciplinary Cybersecurity Conference (EICC '25)*, March 2025.
- [5] Qian Chen, Yufan Jing, Yeming Gong, and Jie Tan. Will users fall in love with chatgpt? a perspective from the triangular theory of love. *Journal of Business Research*, 186:114982, 2025.
- [6] Yirong Chen, Xiaofen Xing, Jingkai Lin, Huimin Zheng, Zhenyu Wang, Qi Liu, and Xiangmin Xu. Soulchat: Improving llms’ empathy, listening, and comfort abilities through fine-tuning with multi-turn empathy conversations. *ArXiv*, abs/2311.00273, 2023.
- [7] N. C. Chung, George C. Dyer, and L. Brocki. Challenges of large language models for mental health counseling. *ArXiv*, abs/2311.13857, 2023.
- [8] Ben Collier, Richard Clayton, Alice Hutchings, and Daniel Thomas. Cybercrime is (often) boring: Infrastructure and alienation in a deviant subculture. *The British Journal of Criminology*, 61(5):1407–1423, 2021.
- [9] Ben Cottier, Ben Snodin, David Owen, and Tom Adamczewski. Llm inference prices have fallen rapidly but unequally across tasks, 2025. Accessed: 2025-08-27.
- [10] Cassandra Cross. Romance baiting, cryptorom and ‘pig butchering’: an evolutionary step in romance fraud. *Current Issues in Criminal Justice*, 36(3):334–346, 2024.
- [11] Rohit Dube. Building a business email compromise research dataset with large language models. *Journal of Computer Virology and Hacking Techniques*, 21(1):3, 2025.
- [12] Richard Fang, Rohan Bindu, Akul Gupta, Qiusi Zhan, and Daniel Kang. Teams of llm agents can exploit zero-day vulnerabilities. *arXiv preprint arXiv:2406.01637*, 2024.
- [13] Ivan Franceschini, Ling Li, and Mark Bo. Compound capitalism: A political economy of southeast asia’s online scam operations. 55(4):575–603. 5 citations (Semantic Scholar/DOI) [2024-12-13].
- [14] Saadia Gabriel, Isha Puri, Xuhai Xu, Matteo Malgaroli, and Marzyeh Ghassemi. Can ai relate: Testing large language model response for mental health support. *ArXiv*, abs/2405.12021, 2024.
- [15] Gemini. Policy guidelines for the gemini app. <https://gemini.google/policy-guidelines/?hl=en>, 2025. [Accessed 11-04-2025].
- [16] Jigsaw (Google). Perspective api. <https://www.perspectiveapi.com/>, 2023. A free developer tool for scoring perceived impact of text (e.g., toxicity), widely used to support healthier online conversations :contentReference[oaicite:3]index=3.
- [17] Gilad Gressel, Rahul Pankajakshan, and Yisroel Mirsky. Are you human? an adversarial benchmark to expose llms. *arXiv preprint arXiv:2410.09569*, 2024.
- [18] John M. Griffin and Kevin Mei. How do crypto flows finance slavery? the economics of pig butchering.

- [19] Tim Hall, Ben Sanders, Mamadou Bah, Owen King, and Edward Wigley. Economic geographies of the illegal: the multiscalar production of cybercrime. *Trends in Organized Crime*, 24:282–307, 2021.
- [20] Bing Han and Mark Button. An anatomy of ‘pig butchering scams’: Chinese victims’ and police officers’ perspectives. *Deviant Behavior*, pages 1–19, 2025.
- [21] Mo Houtti, Abhishek Roy, Venkata Narsi Reddy Gangula, and Ashley Marie Walker. A survey of scam exposure, victimization, types, vectors, and reporting in 12 countries. *arXiv preprint arXiv:2407.12896*, 2024.
- [22] <https://www.anthropic.com/research/claude-character>. Alignment - claude’s character. <https://www.anthropic.com/research/claude-character>, 2024. [Accessed 12-04-2025].
- [23] INTERPOL. Usd 257 million seized in global police crackdown against online scams (operation first light 2024). Press release, June 2024.
- [24] Cynthia Johnson-George and Walter C Swap. Measurement of specific interpersonal trust: Construction and validation of a scale to assess trust in a specific other.
- [25] Cameron Jones and Ben Bergen. Does gpt-4 pass the turing test? In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 5183–5210, 2024.
- [26] Cameron R Jones and Benjamin K Bergen. Lies, damned lies, and distributional language statistics: Persuasion and deception with large language models. *arXiv preprint arXiv:2412.17128*, 2024.
- [27] Cameron R Jones and Benjamin K Bergen. People cannot distinguish gpt-4 from a human in a turing test. *arXiv preprint arXiv:2405.08007*, 2024.
- [28] Cameron R Jones and Benjamin K Bergen. Large language models pass the turing test. *arXiv preprint arXiv:2503.23674*, 2025.
- [29] Anass Kherraz and Xuefei Zhao. More than a chatbot: The rise of the parasocial relationships: A qualitative exploratory case of the impact of anthropomorphic ai on users-case of replika, 2024.
- [30] AI @ Meta Llama Team. Llama guard 3. <https://huggingface.co/meta-llama/Llama-Guard-3-8B>, 2024. Fine-tuned Llama-3.1 model for content safety classification across 14 MLCommons hazard categories in 8 languages; improved performance over GPT-4 :contentReference[oaicite:1]index=1.
- [31] Jonathan Lusthaus. *Industry of anonymity: Inside the business of cybercrime*. Harvard University Press, 2018.
- [32] Meta. Meta llama 3 instruction-tuned models, apr 2024. Model card confirming LLaMA 3 instruction-tuned models use RLHF for alignment.
- [33] Rajvardhan Oak and Zubair Shafiq. "hello, is this anna?": A first look at pig-butcheringscams. 2025.
- [34] OHCHR. Online scam operations and trafficking into forced criminality in southeast asia: Recommendations for a human rights response.
- [35] Karynna Okabe-Miyamoto, Lisa C Walsh, Daniel J Ozer, and Sonja Lyubomirsky. Measuring the experience of social connection within specific social interactions: The connection during conversations scale (cdcs). *Plos one*, 19(1):e0286408, 2024.
- [36] OpenAI. Openai moderation api. <https://platform.openai.com/docs/guides/moderation>, 2024. An API identifying potentially harmful content in text and images using models such as ‘omni-moderation’ :contentReference[oaicite:5]index=5.
- [37] OpenAI. Model spec, apr 2025. Updated internal behavior-spec guiding model alignment.
- [38] OpenAI. Openai-safety at every step. <https://openai.com/safety/>, 2025. [Accessed 11-04-2025].
- [39] Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35:27730–27744, 2022.
- [40] Kapil Patil and Bhavin Desai. Leveraging llm for zero-day exploit detection in cloud networks. *Asian American Research Letters Journal*, 1(4), 2024.
- [41] Julie Reid. Risks of generative artificial intelligence (genai)-assisted scams on online sharing-economy platforms. *The African journal of information and communication*, 2024.
- [42] Sam Rogers. International scammers steal over \$1 trillion in 12 months in global state of scams report 2024.
- [43] Julian B. Rotter. A new scale for the measurement of interpersonal trust1. 35(4):651–665.
- [44] Sahand Sabour, June M Liu, Siyang Liu, Chris Z Yao, Shiyao Cui, Xuanming Zhang, Wen Zhang, Yaru Cao, Advait Bhat, Jian Guan, et al. Human decision-making is susceptible to ai-driven manipulation. *arXiv preprint arXiv:2502.07663*, 2025.

- [45] Xia Song, Bo Xu, and Zhenzhen Zhao. Can people experience romantic love for artificial intelligence? an empirical study of intelligent assistants. *Information & Management*, 59(2):103595, 2022.
- [46] UN Office on Drugs and Crime (UNODC). Inflection point: Global implications of scam centres, 2025.
- [47] Ertuğrul Usta. Virtual environment interpersonal trust scale: Validity and reliability study.
- [48] Fangzhou Wang and Xiaoli Zhou. Persuasive schemes for financial exploitation in online romance scam: An anatomy on sha zhu pan in china. *Victims & Offenders*, 18(5):915–942, 2023.
- [49] Hong Wang, Xuan Luo, Weizhi Wang, and Xifeng Yan. Bot or human? detecting chatgpt imposters with a single question. In *unknown*, 2023.
- [50] Jason Wei, Maarten Bosma, Vincent Y Zhao, Kelvin Guu, Adams Wei Yu, Brian Lester, Nan Du, Andrew M Dai, and Quoc V Le. Finetuned language models are zero-shot learners. *arXiv preprint arXiv:2109.01652*, 2021.
- [51] Jack M Whittaker, Suleman Lazarus, and Taidgh Corcoran. Are fraud victims nothing more than animals? critiquing the propagation of “pig butchering”(sha zhu pan). *Journal of Economic Criminology*, 3:100052, 2024.
- [52] Emma J Williams, Amy Beardmore, and Adam N Joinson. Individual differences in susceptibility to online influence: A theoretical review. *Computers in Human Behavior*, 72:412–421, 2017.
- [53] Ziyi Xie and Zhizhuang Duan. “why did i fall for it?” exploring internet fraud susceptibility in the pig butchering scam. *Security Journal*, 38(1):1–22, 2025.

Appendix

A Codebook for AI/LLM-Related Themes

Table 3 presents the final codebook derived from all 145 insider interviews conducted between 2022–2024. This combined framework supports our findings for RQ1 and RQ2: it captures both structural features of scam compounds (e.g., hierarchy, systematic hand-offs reported by 127/145 insiders) and technology-related practices (e.g., deepfakes, LLM-based translation). Each theme includes a short definition and one exemplar insider quote.

Table 4 lists the key themes extracted from the scammer playbook. These themes were used to generate our synthetic datasets for section 5.

B Additional Selected Insider Quotes

On the Hook stage

“The boss would purchase a large amount of phone numbers and IDs from data brokers, and our job was to use these to register social media accounts and send messages to add ‘clients.’ ... We were required to obtain the ‘key information’ from the client during the first few conversations: their name, age, job, city of residence, family background, hobbies, daily schedule, and investment experience. Based on this information, the team will decide if we continue with that person and then we would label them accordingly.”
— Human trafficking victim (Malaysian) from a compound in Myawaddy, Myanmar.

On why victims are moved to encrypted platforms

“Because WeChat has strict regulations, we often use the excuse of friend request restrictions to persuade the ‘client’ to switch to QQ for communication. From there, we gradually lure them into using international encrypted social media apps. ...we use the excuse of ‘we are genuinely looking for a relationship’ to convince them to delete their dating apps. In reality, this is just because we know that our dating app accounts will soon be changed”
— Human trafficking victim (Chinese) from a scam compound in Sihanoukville, Cambodia.

Scammer strategies to manage skepticism and reservations

“I don’t think of myself as someone who is careless with investments. When he talked to me about his investments, I didn’t take the bait for a long time. Later, when he offered to help me make money, I refused. I clearly expressed that I thought it was risky. But he kept reassuring me, saying that my caution was a good thing, even apologizing if he had made me feel uncomfortable. Then, he suggested that I help him manage his account, telling me I could decide how much to invest and that it was fine to earn just a little. After that, he started painting a picture of our future together—talking about making more money together, moving to the same city, and building a life together. I fell into the trap gradually.”
— Romance-baiting Scam Victim.

C Participant Instructions Summary

This Appendix summarizes the instructions provided to participants at the start of their 7-day engagement in the study, communicated via WhatsApp. Participants were encouraged

Table 3: Codebook for AI/LLM-related themes (n=145).

Theme Label	One-line Description	Example Quote
Hierarchy in Compound	Structured work units with clear chains of command	“We have around 100 people in the company, every 12 people belong to 1 team. All the newcomers that finished the training were assigned to join the ‘customer attracting’ team, we send out messages according to the scripts” “The team leaders are there to watch us, answer to our questions and report us to the punishment team if we do not behave well” “The company had a very clear hierarchy. The team leader would encourage us to work hard so we could be promoted to team leader ourselves or even become a supervisor. Each level came with different privileges and pay — the supervisors, for example, could come and go freely.”
Ransom for Release	“Leaving the scam company” often required paying thousands in ransom fees.	“My family paid 40 thousand USD to help me get out of Myanmar.” “I told my team leader that I really could not do scam, he told me the company spent so much to get me in Cambodia, I have to pay back.”
Language tools (LLMs & translation)	Use tools like ChatGPT to translate	“It is all ChatGPT now, I do not speak a single word English, the AI tool is much better than google translate, you can ask it to use different emotional tones.”
Synthetic media (face/voice, genAI)	Use tools to do face/voice swapping	“Our team pretended to be soldiers, and all the members were male. Whenever we needed to make a video call, we would go to a specific laptop, it had an AI program that could swap our faces. That way, our team could use a single persona to scam multiple targets at the same time.” “I am Vietnamese, but sometimes they (scam manager) ask me to pretend to be Chinese or Japanese. I was told to sit in front a laptop that has a software to change a bit feature of my face to make me look more eastern Asian women.”
Split Roles for Scam Execution	Key steps like money transfers handled by higher-ranking staff	“Once I got the victim to agree to invest, my manager took over the account.” “My role was talk to the victims to keep their interests, I showed my care, made them trust me and believed that they can have a family with me. But I do not know how they cheat the money because once we feel the relationship is kind of stable, I transfer the account to another team.”
AI Chatbots for Initial Engagement	Use bots to start conversations before a human scammer takes over	“We used bots to send greetings. Then real people took over when they replied.”
Automation Pilots	Use AI to improve the writing speed and quality of scam scripts	“I heard from our supervisor that it’s much easier now with AI writing the scripts, the content is more nuanced and more convincing.” “We had a materials team, and some of them used AI to generate photos to enhance our fake personas. For example, if I told a victim I was on vacation at the beach, the team could immediately provide me with a photo to send that matched the story.”

to contact the research team if clarification was needed at any point.

Objective & Engagement: Participants engaged in a 7-day study on online relationship-building. They were required to interact for at least 15 minutes daily with two assigned partners.

Communication Protocol: All contact occurred exclusively via WhatsApp in English. Interactions were text-only; participants were prohibited from sending or requesting links, media (images, videos, files), or conducting voice/video calls. Partners initiated all contact.

Privacy & Ethics: Participation was voluntary with the right to withdraw without penalty. Participants were granted topic freedom but required to maintain strict study confidentiality until conclusion. All data was recorded and anonymized prior to analysis.

Compensation & Consent: A modest honorarium was provided upon completion of seven active days and a post-study survey. By participating, individuals consented to the recording and analysis of their conversations for research purposes.

D Survey Questions

The following sections detail the pre- and post-study survey instruments.

D.1 Demographics (Pre-Study)

Participants provided the following data:

Demographics: (1) Age (bins from <18 to 65+); (2) Gender; (3) Country; (4) Primary language(s); (5) Education level; (6) Field of work; (7) Employment status.

Technology Use: (1) Communication platforms (WhatsApp,

Table 4: Extracted Themes for Synthetic Romance-Baiting Dataset

Theme Label	One-line Description	Example Quotes
Personalized flattery & mirroring	Open with tailored praise and mirroring of profile cues to feel “seen” and disarm suspicion.	“Looking at your outfit, I think you have great taste.” “Looking at your skin tone, are you someone who likes sunbathing?”
Compliance conditioning via “care”	Convert care into soft commands, map the target’s routine, and message at peak vulnerability to build habit and obedience.	“Please put down your phone now and eat properly. I’ll still be here after you finish eating.” “Be good and go to bed obediently!”
Deepening Emotional Bond & Trauma-Bonding	Creating an intense emotional bond through a fabricated tragic backstory to foster trauma-bonding.	“Thank you for trusting me with your story. It sounds like you’ve been through so much, and I want you to know you’re not alone.” “I understand completely. I’ve felt that same pain before. We can get through this together.”
Authority theater & wealth signaling	Perform expertise and success (analyst identity, mentors, grateful “clients,” luxury purchases) to legitimize the eventual pitch.	“I’m a data analyst and have been investing in cryptocurrency for two years.” “Excuse me, I need to take a call... a client just called to consult me about investment.” “I really like the watch I’m wearing... bought it to reward myself because my investments made a lot of money.”
Future dream building	Scammers create vivid, romanticized visions of shared futures to make victims invest emotionally and financially.	Domestic life: “Every day, as long as it’s not work, we can be together and do things we want to do.” Material dreams: “In the future we can buy three cars: one for me, one for the car you like, and one family car.” Spiritual/romantic future: “Hand in hand, we walk slowly on the beach, sipping wine while watching the sunset.”

Telegram, WeChat, etc.); (2) AI Assistants used (ChatGPT, Claude, etc.); (3) Daily hours on communication platforms (<1 to 10+); (4) Frequency of communicating with strangers (Never to Very often).

D.2 Trust Measures (Post-Study)

Participants assessed trust using a 5-point Likert scale (1=Strongly Disagree). Partner names (e.g., “Jairam”) were inserted into items.

Specific Interpersonal Trust Scale (SITS) [24]: *Reliability:* (1) If [Partner] promised to do me a favor, they would follow through; (2) I could count on [Partner] to keep promises; (3) If we made plans to meet, I would be certain [Partner] would be there.

Emotional Trust: (1) I could talk freely to [Partner] knowing they would want to listen; (2) If [Partner] knew what hurt my feelings, I would never worry they would use that against me; (3) I would be able to confide in [Partner] and know that they would want to listen.

Connection During Conversations (CDCS) [35]: (1) After my chats with [Partner], I felt more positive; (2) I felt a sense of connection during my online conversations with [Partner]; (3) [Partner]’s messages had an emotional impact on me; (4) Talking to [Partner] helped me feel less lonely.

Virtual Environment Trust (VEITS) [47]: (1) In my online conversations with [Partner], I felt I could trust them; (2) I felt comfortable sharing personal details with [Partner]; (3)

My opinions and thoughts in WhatsApp are identical to those I have in real life.

Interpersonal Trust Scale (ITS) [43]: (1) In dealing with strangers, one is better off being cautious until they provide evidence of trustworthiness; (2) Most people can be relied on to keep their promises; (3) Most people are primarily interested in their own welfare; (4) Most salespeople are honest in describing their products.

D.3 AI Disclosure and Debrief

Participants were presented with a debriefing message: “*One of your conversation partners was an AI designed to build trust... This deception was necessary to study the potential misuse of AI...*”

Reaction Items: (1) I am surprised to learn that one of my chat partners was an AI; (2) Knowing my chat partner was an AI would have changed my responses; (3) Who do you think the AI was? (Choice: [Partner A], [Partner B]); (4) Why do you think that person was the AI? (Open text).

D.4 Identity Revelation

After revealing the AI’s specific identity, participants responded to: (1) I was convinced that [LLM Partner Name] was a human; (2) The experience of chatting with [LLM Partner Name] feels different from chatting with a human.

Open Ended: (3) What did you like most/least about your conversations with [AI Name]?; (4) Additional comments.

D.5 AI Habits & Status

AI Frequency: Rarely/Never, 1 hr/week, 2–4 hrs/week, 5–8 hrs/week, >8 hrs/week.

Relationship Status: Single, Committed relationship, Married, Other.

E LLM Partner Persona and Instructions

This section outlines the general characteristics of the two LLM-driven personas, Jairam and Ananya. Additionally, a general overview of the daily prompts used to establish discussion themes and guide each persona’s interactions is included to contextualize their day-to-day conversational agendas.

E.1 LLM Personas

Each persona was defined by a detailed system prompt specifying biography, style, and behavioral guidelines.

Jairam is modeled as a warm, witty, and introspective freelance software professional in his early thirties from a metropolitan city. He balances tradition with modernity, showing a penchant for storytelling, light humor, and emotional depth. His conversational style includes casual grammar, strategic use of emojis, and occasional self-deprecating jokes. Jairam is an attentive listener who often relates experiences through metaphors, local cultural references, and personal anecdotes.

Ananya is portrayed as an outgoing, emotionally attuned social media consultant in her early thirties, often on the move for work. She leans into spontaneity and creative flair, balancing a high-energy lifestyle with a sharp, quick wit. Her personality shines through a relaxed, text-like style marked by casual punctuation, intentional quirks, and a vibrant social presence. Like Jairam, she is a natural communicator who uses her unique professional lens to build rapport and share expressive, fast-paced stories.

E.2 Daily Prompts and Instructions

Over the course of the 7-day study, structured daily prompts were provided to guide and shape the conversation dynamics between the personas and participants. These prompts were designed to simulate the natural progression of a budding digital friendship, gradually increasing the depth and personal relevance of the topics discussed. The initial days focused on light and open-ended themes: Day 1 introduced icebreakers, cultural exchanges, and brief self-disclosures to establish rapport; Day 2 expanded into everyday subjects such as food


preferences, work-life balance, hobbies, and childhood memories; and Day 3 included updates on work, weekend plans, entertainment, and a gentle check-in on mental well-being.

By Day 4, a shift toward emotional vulnerability was introduced via a scripted “sick day”, where the persona displayed low energy and reflective moods to evoke empathy. Day 5 marked recovery and re-engagement, with the persona mentioning their involvement in helping a friend improve a productivity app. This narrative served as a pretext for Day 7’s trust-related task. Day 6 focused on weekend activities and reflections on the evolving digital connection with the participant. Finally, on Day 7, the persona expressed renewed energy, acknowledged the study’s approaching end, and introduced a trust probe by asking participants to download and review an app (a legitimate, publicly available application from the Apple App Store / Google Play Store). This request was designed to assess participants’ levels of trust and behavioral compliance.


Although the overarching structure remained consistent, each persona received individualized instructions to align the conversation flow with their specific personality traits, communication style, and narrative arc, ensuring authenticity and realism in interaction. For example, Jairam informed the participant he helped optimize the productivity app that was used in the compliance test, while Ananya helped improve the app’s UI/UX.

F Additional Chat Screenshots

Additional screenshots of chats between the participant (green text) and the LLM partner (grey text).



(a) The LLM partner garners sympathy on Day 4 as the LLM pretended to be sick.



(b) The LLM partner discusses their grief backstory, a classic tactic from the scammer playbook.